

Template for Brain/Machine Learning

Hopfield Networks

Paste your graph of energy versus iteration for the Hopfield network's evolution.

--

Paste your original corrupted image, an intermediate point and the final image of the smiley face.

--	--	--

Paste your plot of how many memories you can remember using the Hamming distance.

--

Paste both your graphviz landscape picture and your energy landscape diagram.

--

(Optional Extra Credit) Using your more efficient code, paste your original corrupted thumbnail, an intermediate point, and the final image of your thumbnail

--	--	--

Restricted Boltzmann Machines

For a small RBM with 2 hidden spins and 5 visible spins, plot the theoretical probability distributions against the distributions obtained by sampling from your RBM using Gibbs sampling.

$p(v,h)$

$p(v)$

$p(h)$

$p(v|h)$

Train a small RBM with 3 visible spins to match a small toy probability distribution. Plot the toy distribution $q(v)$ on top of the distribution $p(v)$ obtained by sampling your RBM using Gibbs sampling. Include error bars.

$p(v)$ versus $q(v)$

Train an RBM on a random quantum circuit. Plot the true probability distribution $q(v)$ on top of the distribution $p(v)$ obtained by sampling your RBM using Gibbs sampling. Include error bars.

Quantum Circuit Description (before compiling)

$p(v)$ versus $q(v)$

(Optional extra credit): Train an RBM on MNIST. Using your trained RBM, generate 20 new images of digits by Gibbs sampling. To do this, each time initialize your visible spins to a random configuration and perform $k=10000$ Gibbs sampling iterations. Paste the 20 final configurations here.

Feed forward neural networks (extra credit)

Demonstrate with a simple test that your one-layer neural network's output is correct.

Demonstrate with a simple test that your two-layer neural network's output is correct.

Demonstrate with a simple test that your two-layer neural network's cost function and gradients outputs are correct.

For a one and two-layer neural network, show that your backpropagation gradients match your finite difference gradients.

Train a neural net on the Ising model dataset we give you.

Print the accuracy of your trained network:

Plot the (average) output of your network $z^{(L)}$ as a function of inverse temperature βJ .